

Efficient coinductives through state-machine corecursors

Supervisor: Alex Keizer

Marking supervisor: Jamie Vicary

Director of studies: Russell Moore

Word count: 951

1 Work to be undertaken

Previous work has shown how coinductive types can be encoded in Lean as quotients of polynomial functors (QPFs). This encoding is expressive but inefficient for cofixed points. For example accessing the n index of a stream takes $\mathcal{O}(n^2)$ time. This becomes an issue when using Lean as a general purpose programming language. An alternative approach is using a state-machine based encoding of cofixed points. This representation directly encodes the coalgebraic structure and is well understood. The goal of this project is formalising the equivalence between these two representations. This equivalence should come directly from the corecursors of each of the implementations. With these we get nice computational behaviour of the compiled code, meaning directly getting a stream to depth n is $\mathcal{O}(n)$. In Figure 1 we can see in blue the performance from the current implementation, versus in orange the performance of a prototype implementation the state-machine based version.




Figure 1: Graph plotting current performance v possible gains

2 Core project

This project will be implementing the state-machine representation of coinductive types. The goal will be formalising the equivalence between this and the progressive approximation representation. I will start with the equivalence between `Streams` in the two representations, then building up to the general cofix structure.

After this I will be implementing a coinductive monad allowing us to encode non-termination as an effect using the state-machine representation. This will simply be an example demonstrating the first part of the project.

3 Starting point

I have worked with QPFs on the meta-programming side for an internship between my Part Ia and Part Ib. During this I learnt of the basics of polynomial functors, and `TypeVecs`. This means I am aware of what the underlying structures are when it comes to the raw implementation. I have also done a feasibility assessment of the project by seeing how the current polynomials respond to universe levels. This led to me making 2 pull requests (#28095, #28279) to mathlib on `TypeVec` in preparation for my project. I also tried to merge (#28112) an implementation of commutable `Shrink` to mathlib that later got reverted when it was found to be inconsistent.

There are more minor refactoring pull requests towards the mathlib repository which don't change any behaviour but in general all of these can be *found on this link*. I include these for completeness and transparency.

4 Substance

4.1 The \mathbb{M} type

The \mathbb{M} type is the name given to the terminal coalgebra of polynomial functors; the possibly infinitely deep trees. These are generated by progressive approximation where earlier trees must “agree” with the later ones. Agreement is given by them being the same up to the previous depth. A visual example is given in Figure 2. We can have approximations for any n , thereby letting the trees take any depth including infinite depth. Trees can be terminated by having no children.




Figure 2: Agreement of two trees of height 1 and 2

The current QPF implementation has 2 \mathbb{M} types; univariate and multivariate implementations. These both have the undesired computational behaviour. \mathbb{M} types are polynomial.

4.2 The `Cofix` type

`Cofix` is the terminal coalgebra in QPFs, (possibly) infinitely big quotiented trees. This is the slightly concerning part of the project and by far the highest risk section as working with `Quot` in Lean is a painful experience.

4.3 The non-termination monad


The non-termination monad is a simple example of a coinductive. This will be useful for testing the state-machine representation's performance, I will be implementing this. The structure of the monad as a coinductive has two constructors as seen in this pseudocode:

```
coinductive NTMonad (A : Type)
| val : A -> NTMonad A
| tau : NTMonad A -> NTMonad A
```

5 Evaluation

The success of this project can be given by how close to the theorised performance we can get to. The goal would be getting to the same order or magnitude.

6 Core timeline



The plan for work would be divided into a few different stages.

6.1 Variable universe Ms (2025-10-11 2w 2025-10-24)

To begin with I have to make the corecursor universe polymorphic. As stated in Section 3 I have most of this code written. I will attempt to get this merged into mathlib, but if this review process takes too long I will be working on my own branch.

6.2 Stream special example (2025-10-25 2w 2025-11-7)

The first step would be implementing `Stream` under the two representations. Then I will familiarise myself with using bisimilarity to formalise the equivalence. This is a first step to just understand how these features work.

6.3 Univariate M (2025-11-8 2w 2025-11-21 + (CAT exam) + 2025-12-05 3w 2025-12-23)

Implementing the univariate `M` type is the natural next step from this. This will be the natural next step in difficulty. The work should lay the groundwork for the next proofs.

6.4 Multivariate \mathbf{M} (2025-12-26 4w 2026-01-23)

The next step will be generalising the univariate implementation, the same structure of the proof should be the same for the multivariate case. The main difficulty here comes from working with `TypeVecs`. These are quite difficult to reason about.

6.5 `Cofix` (2026-01-24 6w 2026-03-06)

Finally, the `Cofix` type has to be implemented. This will be the most challenging part as I will have to work with quotients.

6.6 Implementing the `NTMonad` (2026-03-07 2w 2026-03-20)

This will be a demonstration of all of the previous work.

7 Extensions

7.1 Shrinking the representations (2026-03-21 2w 2026-04-04)

This involves finding a sound way allowing the current theory to use the efficient representation. The soundness of this will be difficult to reason about as it requires working with the Lean code generator. Therefore I leave this as an extension.

8 Resources

Access to the restricted side of the lab is needed for working with the greater team.